

Architecting Modern Data Pipelines for Sports Analytics: A Comparative Analysis of AWS and Databricks Lakebase Systems

Introduction to the Analytical Paradigm

The proliferation of high-fidelity data collection mechanisms in professional football (soccer) has catalyzed a fundamental paradigm shift in sports analytics. The transition from rudimentary, manual statistical counting to advanced spatio-temporal modeling necessitates highly robust data engineering infrastructures. Modern football analytics relies heavily on massive datasets, ranging from high-frequency optical tracking data capturing player coordinates multiple times per second to granular, manually annotated event logs detailing every on-ball action.¹ To extract actionable intelligence from these disparate data streams, organizations require sophisticated data pipelines capable of ingesting, transforming, and orchestrating analytical workflows.

This comprehensive report provides a strategic architectural blueprint for building a production-grade data system tailored to support advanced sports analytics curricula, specifically addressing the rigorous requirements of the Soccermetrics Pro cohort.³ The analysis deconstructs a specific target architecture provided in the initial system schematic, mapping its components to a traditional Amazon Web Services (AWS) data warehousing paradigm. This involves orchestrating Amazon Simple Storage Service (S3), AWS Glue, Amazon Redshift, and PostgreSQL using Apache Airflow and the data build tool (dbt).⁴ Subsequently, the report contrasts this traditional pipeline with a modernized implementation utilizing the newly released Databricks Lakebase service, a serverless PostgreSQL-compatible database integrated directly into the lakehouse architecture.⁷ To facilitate the practical implementation of these systems, the report also curates a definitive catalog of confirmed open-source soccer datasets and identifies critical AWS Skill Builder learning modules.⁹

Data Foundations for Advanced Football Analytics

The Soccermetrics Pro curriculum operates at a late-undergraduate to master's academic level, specifically targeting data scientists, match analysts, and technical personnel within professional football organizations.³ The mathematical and statistical rigor of the course—encompassing logistic regression, Markov chains for Expected Threat (xT), Voronoi diagrams for spatial pitch control, and various machine learning methodologies—demands a

highly structured and reliable data foundation.³

Confirmed Open-Source Data Ecosystems

To build a realistic pipeline that accurately simulates the proprietary environments of elite football clubs, engineers must source high-fidelity, multidimensional data. The following open-source datasets are confirmed as core pillars of the Soccermatics curriculum and provide the necessary structural complexity for pipeline engineering.²

Data Provider	Data Type	Structural Format	Key Characteristics and Analytical Utility
StatsBomb	Event & 360 Context	Nested JSON	Highly granular event logs (over 3,400 events per match). Includes specialized metrics such as pressure events, pass heights, and shot freeze-frames. Accessible via the statsbombpy API or the statsbomb/open-data GitHub repository. ⁹
Metrica Sports	Optical Tracking	CSV / JSON	High-frequency spatial data capturing the x and y coordinates of all 22 players and the ball at 25 frames per second. Critical for pitch control models and calculating player

			velocity. ²
Wyscout	Global Event Data	JSON	Expansive event datasets covering major global competitions, including the top five European leagues. Provides massive sample sizes necessary for training robust machine learning models like Expected Goals (xG). ¹⁵

The structural hierarchy of the StatsBomb dataset, for instance, presents a specific engineering challenge for the transformation layer. The data is organized around a relational indexing chain: competitions.json contains competition and season metadata, which maps to matches/{competition_id}/{season_id}.json. These match files subsequently link to highly nested event payloads stored in events/{match_id}.json and lineups/{match_id}.json.²⁰ Flattening these nested JSON structures into columnar formats suitable for analytical querying is a primary objective of the ETL pipeline.

Target Analytical Tooling

The data pipeline must be engineered to deliver optimized, query-ready data to downstream data scientists operating in Python environments. The Soccermatics syllabus relies heavily on specific Python libraries that must interface seamlessly with the underlying database or data warehouse. The mplsoccer library is utilized extensively for rendering spatial visualizations, such as pass networks, shot maps, and heat maps.³ For application deployment, the curriculum mandates the use of Streamlit to construct interactive web applications that allow end-users (e.g., coaches or scouting directors) to interface with the analytical models.³ The requirement to serve live web applications dictates that the final serving layer of the database architecture must support high-concurrency, low-latency queries.

Deconstruction of the Target System Architecture

The architectural schematic provided as the target reference model depicts a highly structured, sequential data flow transitioning from raw ingestion to a final operational database layer. Analyzing the specific nodes of this visual representation is critical for establishing the

baseline implementation.

The pipeline initiates with "Data Providers" feeding into a "DataLab VM AWS EC2" instance hosting "Automated Apache Airflow." This indicates a self-managed orchestration environment where Airflow schedulers and workers operate directly on Elastic Compute Cloud (EC2) virtual machines. Airflow executes a "WRITE" operation, landing the raw data payloads into a "Datalake AWS S3" bucket. Following ingestion, an "EXTRACT" operation is triggered, utilizing "ETL AWS Glue" to process the raw storage. The transformed data is subjected to a "LOAD" operation into the "Data Warehouse AWS Redshift."

Within the Redshift environment, the architecture highlights a cyclical transformation process managed by "dbt" (data build tool). This represents the modern Extract, Load, Transform (ELT) paradigm, where data is loaded into the warehouse before complex SQL-based business logic is applied to model the data into facts and dimensions. Finally, the schematic depicts a secondary "LOAD" operation transferring the modeled data out of Redshift and into an "OLTP DB Postgres" instance.

This specific architectural pattern—moving data from an Online Analytical Processing (OLAP) warehouse into an Online Transactional Processing (OLTP) database—is known as Reverse ETL. It is implemented because columnar data warehouses like Redshift are optimized for massive, complex aggregations but perform poorly when serving the high-volume, single-row, low-latency queries required by interactive web applications like Streamlit.²² The PostgreSQL database acts as the high-speed serving layer for the end-user application.

Implementation I: The Traditional AWS Data Warehouse Pipeline

Replicating the target architecture requires orchestrating a sequence of decoupled AWS services. This traditional implementation offers high flexibility and leverages mature, battle-tested components, making it an ideal pedagogical environment for mastering core data engineering principles.²⁵

Orchestration and Ingestion Operations

While the schematic depicts Apache Airflow hosted on an EC2 instance, enterprise best practices dictate transitioning to Amazon Managed Workflows for Apache Airflow (MWAA). Managing Airflow on raw EC2 instances introduces significant operational overhead, including provisioning infrastructure, managing auto-scaling for worker nodes, and securing the environment. Amazon MWAA provides a secure, fully managed orchestration service that automatically scales the underlying containerized workers to match the workload demands of the pipeline.⁶

The ingestion phase is codified within an Airflow Directed Acyclic Graph (DAG) using Python.

The DAG initiates Python tasks utilizing the statsbombpy library and standard API request modules to fetch the JSON event data and CSV tracking data from the respective GitHub repositories.²⁸ The Airflow workers subsequently write these raw, unmodified files directly into the raw zone of the Amazon S3 data lake. To transition from a batch-polling model to a highly responsive event-driven architecture, S3 Event Notifications can be configured. When a new match file lands in the S3 bucket, it triggers an AWS Lambda function that alerts the Airflow API to commence the transformation DAG, significantly reducing pipeline latency.²⁵

Serverless Data Integration via AWS Glue

Because sports data—particularly the StatsBomb event payloads—features deeply nested arrays and dictionaries (e.g., a pass event containing nested dictionaries for pass length, angle, recipient, and pressure context), it is highly inefficient to query in its raw JSON state. AWS Glue serves as the serverless data integration engine to sanitize and normalize this data.¹⁰

The process begins with AWS Glue Crawlers, which automatically scan the S3 raw zone to infer the schema of the incoming football data and populate the AWS Glue Data Catalog.³¹ This metadata catalog provides a centralized schema registry. Subsequently, AWS Glue ETL jobs, executing distributed Apache Spark code, extract the JSON data. The Spark transformations flatten the nested structures into tabular rows. For instance, a single event containing an array of spatial coordinates is exploded into distinct columns for x_{start} , y_{start} , x_{end} , and y_{end} .

Crucially, the AWS Glue job converts the JSON and CSV files into Apache Parquet, a columnar storage format.³³ Parquet applies aggressive compression and encoding, which dramatically reduces the amount of data scanned during analytical queries, optimizing both performance and cost when the data is eventually queried.

Warehousing and Transformation with Amazon Redshift and dbt

The transformed Parquet files are loaded into Amazon Redshift. Redshift is a Massively Parallel Processing (MPP) data warehouse that distributes query execution across multiple compute nodes, making it highly proficient at running complex aggregations over millions of rows of historical match data.³⁴

Loading data from S3 into Redshift is executed using the COPY command, which allows Redshift to read and load data in parallel directly from the S3 buckets, bypassing the Airflow worker memory.³⁵ Once the raw tables exist in Redshift, the pipeline delegates the data modeling responsibility to dbt. By connecting dbt to Redshift via the dbt-redshift adapter, analytics engineers can define modular data transformations using standard SQL SELECT statements.⁴

In the context of the Soccermetrics curriculum, dbt is utilized to build the analytical models required for advanced metrics. A dbt model might take the raw event data and execute the

spatial geometry equations necessary to calculate the distance to the goal for every shot, creating a new "feature mart" table. dbt introduces software engineering rigor to SQL, allowing developers to implement version control, automated testing (such as ensuring a `match_id` is never null), and automated dependency mapping.⁴

To optimize Redshift for the specific querying patterns of spatial tracking data, dbt configuration blocks must specify distribution keys (`distkey`) and sort keys (`sortkey`). For example, configuring the `match_id` as the distribution key and the timestamp as the sort key ensures that queries calculating player trajectories over a specific time window are executed with maximum efficiency, as the relevant data blocks are co-located on disk.³⁸

The Reverse ETL Implementation

The final stage of the target architecture requires pushing the transformed analytical data from Redshift into a PostgreSQL database. While Redshift excels at scanning billions of rows to calculate seasonal averages, it suffers from high latency when executing thousands of concurrent, single-row lookup queries—the exact pattern generated by an interactive Streamlit web application visualizing a specific pass network.²²

To populate the operational PostgreSQL database (such as Amazon RDS), the pipeline implements a Reverse ETL pattern. An Airflow operator triggers a Redshift UNLOAD command, which writes the finalized dbt models back to an S3 bucket. A subsequent Airflow task utilizes PostgreSQL's `s3_import` functionality to ingest the data into the operational database.²³ This dual-database architecture ensures that the analytical heavy lifting is isolated in the warehouse, while the frontend application benefits from the sub-millisecond response times of the OLTP system. However, this pattern introduces significant architectural complexity, data duplication, and maintenance overhead.

Implementation II: Databricks Lakebase Architecture

The traditional architecture, while robust, highlights a historical dichotomy in data engineering: the rigid separation between analytical warehouses (Redshift) and operational databases (PostgreSQL). Databricks Lakebase, introduced to general availability in February 2026, represents a fundamental architectural evolution designed to collapse this dichotomy, offering a streamlined alternative for the Soccermatics pipeline.⁸

The Decoupled Database Paradigm

Databricks Lakebase is a fully managed, serverless OLTP database compatible with PostgreSQL 17. It is defined by a "lakebase" architectural category that entirely separates compute resources from storage.⁸ In a traditional operational database, storage and compute are tightly coupled on the same underlying hardware. Consequently, if a data scientist runs a heavy analytical query (such as training an Expected Goals model on five years of data), it consumes CPU and memory resources, potentially causing latency spikes or timeouts for live users

interacting with the Streamlit application.⁴⁰

Lakebase mitigates this resource contention by utilizing lightweight, ephemeral compute clusters that operate on top of durable data lake storage.⁸ The Autoscaling version of Lakebase automatically provisions and scales compute capacity based on instantaneous workload demands. During periods of inactivity, the compute can scale to zero, ensuring that organizations only pay for active processing time, calculated in Databricks Units (DBUs).⁸

Zero-ETL and the Unified Lakehouse

Implementing the Soccermetrics pipeline on Databricks eliminates the need for the complex AWS Glue and Redshift COPY orchestration. Raw JSON and CSV data from StatsBomb and Metrika are ingested directly into Delta Lake tables residing in cloud object storage (the foundational layer of the Databricks Data Intelligence Platform). Delta Lake provides ACID transaction guarantees and schema enforcement to the raw data files.⁴³

Instead of relying on fragile Reverse ETL pipelines to move data from the warehouse to the serving database, Lakebase implements a "Zero-ETL" philosophy through Synced Tables. Leveraging managed Lakeflow Spark Declarative Pipelines, the curated, modeled data residing in Delta Lake is continuously and asynchronously synchronized into the Lakebase PostgreSQL interface.⁴⁵

Because Lakebase writes directly to the underlying lakehouse storage, analytical engines like Apache Spark and Databricks SQL can query the operational data immediately without requiring data extraction.⁸ This unified environment allows data scientists to build complex pitch control models using Spark's distributed computing capabilities, while the Streamlit application simultaneously queries the exact same unified dataset via standard PostgreSQL JDBC drivers with sub-10ms latency.⁴⁷

Advanced Developer Workflows and AI Integration

The Lakebase architecture introduces advanced developer capabilities that profoundly benefit the rapid experimentation required in an academic or professional analytics cohort. Because compute and storage are decoupled, Lakebase supports instant database branching utilizing copy-on-write technology.⁷ A data scientist can instantly clone the entire production database, including schemas and data, to test a new Markov chain model for Expected Threat. Destructive tests can be executed on this isolated branch without impacting the live application, and the branch can be discarded or merged as needed.⁷ Furthermore, automated snapshots provide point-in-time recovery, allowing teams to roll back the database to a precise state in the event of an erroneous transformation.⁵⁰

The modeling phase remains rigorous through the use of dbt. Rather than utilizing the dbt-redshift adapter, the Databricks implementation utilizes the dbt-postgres adapter to

execute transformations directly within the Lakebase environment.⁵² Furthermore, because Lakebase supports standard PostgreSQL extensions, it natively includes pgvector.⁸ This allows the pipeline to store and query high-dimensional vector embeddings, a critical capability for integrating Large Language Models (LLMs) or executing similarity searches on player movement patterns, directly within the operational database.⁸

Comparative Architectural Analysis

Evaluating the traditional AWS architecture against the Databricks Lakebase paradigm reveals distinct strategic trade-offs, particularly in the context of supporting the multifaceted workloads of the Soccermatics Pro curriculum.

Architectural Component	Traditional AWS Implementation	Databricks Lakebase Implementation
Primary System Archetype	Massively Parallel Processing (MPP) OLAP Data Warehouse feeding an OLTP replica ²³	Unified OLTP Database directly integrated with Lakehouse storage ⁸
Data Storage Format	Proprietary Redshift columnar format (requires data movement from S3) ⁵⁵	Open formats (Delta Lake, Apache Parquet) natively queried ⁵⁶
Transformation Tooling	AWS Glue (Spark) for normalization, dbt-redshift for modeling ³⁷	Databricks Spark for normalization, dbt-postgres for modeling ⁵²
Application Serving Layer	Requires Reverse ETL via S3 UNLOAD to Amazon RDS PostgreSQL ²³	Natively serves applications with low latency (<10ms) via Synced Tables ⁴⁷
Orchestration Overhead	High: Requires manual DAG configuration in Amazon MWAA to coordinate disparate services ⁴	Low: Handled automatically via Unity Catalog and Lakeflow Pipelines ⁴⁵
Compute Scaling Mechanism	Traditionally clustered compute scaling (Serverless options exist but distinct) ⁴⁰	Independent ephemeral compute; scales to zero automatically ⁴⁰

Developer Environments	Requires manual snapshotting and parallel cluster provisioning ⁵⁸	Instant zero-copy database branching and point-in-time recovery ⁷
Vector and AI Integration	Requires external integrations (e.g., Amazon SageMaker) ²⁴	Native integration via pgvector within PostgreSQL 17 ⁸

The traditional AWS architecture remains highly performant for organizations that prioritize rigid, structured reporting and massive historical aggregations where Redshift's vectorized querying excels.³⁴ However, the operational overhead required to manage the physical infrastructure—including IAM roles, VPC configurations, and complex Airflow DAGs across S3, Glue, Redshift, and RDS—diverts significant engineering resources away from actual data science.⁴ The Reverse ETL pattern introduces a fragile dependency chain where application data is inherently latent.

Conversely, Databricks Lakebase provides a superior paradigm for the agile, machine-learning-heavy workloads characteristic of modern sports analytics. By unifying the analytical and operational planes, Lakebase eliminates redundant data movement.⁸ The ability to execute complex Spark-based physics models on tracking data and immediately serve the modeled outputs to a Streamlit application via PostgreSQL protocols drastically accelerates the time-to-insight. The inclusion of instant branching and scale-to-zero compute provides a highly cost-effective and safe sandbox environment, which is highly advantageous for an academic or professional training cohort.⁷

Prescriptive AWS Skill Builder Learning Plan

For practitioners tasked with mastering the traditional AWS architecture depicted in the target schematic, leveraging official training resources is critical. Utilizing an active AWS subscription, the following confirmed AWS Skill Builder modules and workshops constitute a comprehensive learning plan to acquire the practical experiences necessary to implement the S3, Glue, MWAA, and Redshift pipeline.⁵⁹

1. Data Integration and Serverless ETL

Mastering the transformation of nested JSON football data requires proficiency in AWS Glue.

- **AWS Glue Getting Started (Digital Course):** This foundational module introduces the theoretical concepts of serverless data integration. It covers the deployment of the AWS Glue Data Catalog, which is essential for schema inference on the StatsBomb datasets, and the utilization of AWS Glue Studio for visual job creation.¹⁰
- **Data Engineering on AWS - A Data Lake Solution (Digital Course with Labs):** A comprehensive four-hour curriculum that provides hands-on builder labs. It explicitly covers the architectural patterns for building resilient S3 data lakes and utilizing AWS Glue

to execute the extraction and loading phases.⁶¹

2. Data Warehousing Optimization

Amazon Redshift requires specific tuning to efficiently query massive spatio-temporal datasets.

- **Amazon Redshift Getting Started (Digital Course):** A one-hour primer on provisioning clusters, managing namespaces, and executing basic query operations.⁶²
- **Amazon Redshift – Best Practices for Data Warehousing on AWS:** This critical intermediate course dictates the strategies for advanced data modeling. It provides deep instruction on configuring column compression, selecting optimal distribution keys (distkey), and defining sort keys (sortkey). Mastering these concepts is mandatory for optimizing the analytical queries that process high-frequency optical tracking data.⁶³
- **Working with Amazon Redshift (AWS Builder Lab):** A two-hour, subscription-based practical lab that provides hands-on experience executing the COPY command to load raw data files in parallel from Amazon S3 directly into the Redshift compute nodes.⁶²

3. Orchestration and Transformation Engineering

To automate the pipeline and integrate the data build tool (dbt), advanced orchestration skills are required.

- **Amazon Managed Workflows for Apache Airflow (MWAA) Getting Started:** A fundamental course that guides users through the AWS Management Console to provision secure MWAA environments and deploy rudimentary Python DAGs to automate data flows.¹¹
- **Transforming Data in S3 with Redshift & dbt (AWS Workshop Studio):** A highly specialized Level-300 workshop. This module provides step-by-step instruction on connecting an open-source dbt project to an Amazon Redshift cluster. It covers the authoring of SQL transformation models, the implementation of automated data quality tests, and the generation of standardized documentation.⁶⁵
- **Amazon MWAA for Analytics Workshop:** An advanced, two-hour lab dedicated to building complex, cross-service data pipelines. It provides exposure to the specific Airflow operators and hooks required to trigger AWS Glue jobs and Redshift queries systematically.⁶⁸

4. Holistic Certification Pathway

For a comprehensive validation of these competencies, the **AWS Certified Data Engineer - Associate (DEA-C01) Learning Plan** aggregates these concepts. This twenty-hour curriculum covers the orchestration of data pipelines, troubleshooting serverless analytics, and managing the entire data lifecycle across S3, Glue, and Redshift, ensuring a profound mastery of the traditional architecture.⁶¹

Conclusion

The architecture required to support advanced sports analytics—encompassing Expected Threat modeling, pitch control calculations, and dynamic spatial visualizations—must be highly resilient and deeply integrated. The traditional AWS pipeline, characterized by the sequential orchestration of S3, Glue, Redshift, and PostgreSQL via MWAA and dbt, offers a highly modular, strictly governed, and deeply mature ecosystem. Mastery of this architecture provides profound expertise in distributed computing and enterprise data engineering.

However, the comparative analysis indicates that the Databricks Lakebase architecture resolves the fundamental friction points inherent in traditional designs. By decoupling compute from storage and uniting the OLAP lakehouse with a serverless OLTP PostgreSQL interface, Lakebase eradicates the need for fragile Reverse ETL processes. It allows data scientists within the Soccermatics Pro cohort to query raw optical tracking data, execute complex mathematical models, and serve those metrics to interactive web applications through a singular, unified platform. This structural elegance minimizes infrastructure overhead, enabling analysts to dedicate their resources to the extraction of tactical intelligence from the beautiful game.

Works cited

1. Why More Teams Are Switching to AI-Powered Sports Analytics, accessed February 25, 2026, <https://wsc-sports.com/blog/industry-insights/why-more-teams-are-switching-to-ai-powered-sports-analytics/>
2. Team Sports Analytics - DTAI, accessed February 25, 2026, <https://dtai.cs.kuleuven.be/tutorials/sports/ecml2024/notes/datasets/soccer>
3. Welcome to Soccermatics — Soccermatics documentation, accessed February 25, 2026, <https://soccermatics.readthedocs.io/en/latest/index.html>
4. Build data pipelines with dbt in Amazon Redshift using Amazon MWAA and Cosmos, accessed February 25, 2026, <https://aws.amazon.com/blogs/big-data/build-data-pipelines-with-dbt-in-amazon-redshift-using-amazon-mwaa-and-cosmos/>
5. Build a Data Warehouse using AWS Glue, S3 & Amazon Redshift - YouTube, accessed February 25, 2026, <https://www.youtube.com/watch?v=ww2dhCYqJNl>
6. Amazon Managed Workflows for Apache Airflow (MWAA), accessed February 25, 2026, <https://aws.amazon.com/managed-workflows-for-apache-airflow/>
7. Lakebase - Databricks, accessed February 25, 2026, <https://www.databricks.com/product/lakebase>
8. Databricks Introduces Lakebase, a PostgreSQL Database for AI Workloads - InfoQ, accessed February 25, 2026, <https://www.infoq.com/news/2026/02/databricks-lakebase-postgresql/>
9. statsbomb/open-data: Free football data from StatsBomb - GitHub, accessed February 25, 2026, <https://github.com/statsbomb/open-data>

10. AWS Glue Getting Started - Request Server Error Page, accessed February 25, 2026,
<https://skillbuilder.aws/learn/FJ1W4H4Y4J/aws-glue-getting-started/XTFXU583XT>
11. Amazon Managed Workflows for Apache Airflow (MWAA) Getting Started - AWS Skill Builder, accessed February 25, 2026,
<https://skillbuilder.aws/learn/QMFVM6GAAU/amazon-managed-workflows-for-apache-airflow-mwaa-getting-started/YNECTHRZ93>
12. Soccermatics Pro - Twelve Football, accessed February 25, 2026,
<https://twelve.football/courses/soccermatics-pro>
13. OpenSTARLab: Open Approach for Spatio-Temporal Agent Data Analysis in Soccer - arXiv, accessed February 25, 2026, <https://arxiv.org/html/2502.02785v2>
14. Hudl Statsbomb - The World's Most Advanced Soccer Data, accessed February 25, 2026, <https://www.hudl.com/products/statsbomb>
15. 20 Best Free Sports Datasets for ML 2025 - Unidata, accessed February 25, 2026,
<https://unidata.pro/blog/best-free-sports-datasets-ml/>
16. Metrica Sports sample tracking and event data - GitHub, accessed February 25, 2026, <https://github.com/metrica-sports/sample-data>
17. Friends-of-Tracking-Data-FoTD/LaurieOnTracking - GitHub, accessed February 25, 2026, <https://github.com/Friends-of-Tracking-Data-FoTD/LaurieOnTracking>
18. koenvo/wyscout-soccer-match-event-dataset - GitHub, accessed February 25, 2026, <https://github.com/koenvo/wyscout-soccer-match-event-dataset>
19. Wyscout - kloppy 3.18.0 - PySport, accessed February 25, 2026,
<https://kloppy.pysport.org/reference/providers/wyscout/>
20. Complete guide on working with the StatsBomb Open Data dataset (w/ applications), accessed February 25, 2026,
<https://medium.com/@lucascarrasquillaparra/complete-guide-on-working-with-the-statsbomb-open-data-dataset-a57c26d5852b>
21. Getting Started — Soccermatics documentation, accessed February 25, 2026,
<https://soccermatics.readthedocs.io/en/latest/gettingstarted.html>
22. Redshift vs Databricks | Performance & Pricing: Comparison Guide - Firebolt, accessed February 25, 2026,
<https://www.firebolt.io/comparison/redshift-vs-databricks>
23. How We Built 15Five's Data Platform, accessed February 25, 2026,
<https://www.15five.com/blog/how-we-built-15fives-data-platform>
24. AWS Redshift Architecture: 5 Important Components - Airbyte, accessed February 25, 2026,
<https://airbyte.com/data-engineering-resources/aws-redshift-architecture>
25. Create a modern data platform using the Data Build Tool (dbt) in the AWS Cloud, accessed February 25, 2026,
<https://aws.amazon.com/blogs/big-data/create-a-modern-data-platform-using-the-data-build-tool-dbt-in-the-aws-cloud/>
26. AWS serverless data analytics pipeline reference architecture | AWS Big Data Blog, accessed February 25, 2026,
<https://aws.amazon.com/blogs/big-data/aws-serverless-data-analytics-pipeline-reference-architecture/>

27. Best practices for least privilege configuration in Amazon MWAA | AWS Big Data Blog, accessed February 25, 2026,
<https://aws.amazon.com/blogs/big-data/best-practices-for-least-privilege-configuration-in-amazon-mwaa/>
28. Use Statsbombpy with Python - Hudl Support, accessed February 25, 2026,
https://support.hudl.com/s/article/use-statsbombpy-python?topic=Statsbomb_Python_Tools
29. Load Data Into an S3 Bucket in Amazon Web Services Using Python - YouTube, accessed February 25, 2026, <https://www.youtube.com/watch?v=IO5CyF5q5MM>
30. Airflow Event-Driven Architecture: Automate S3 → Lambda → Redshift → SageMaker Pipeline - YouTube, accessed February 25, 2026,
<https://www.youtube.com/watch?v=7CAXKJiVTYs>
31. How I Prepared for the AWS Data Engineer Associate Exam (And What I'd Do Differently), accessed February 25, 2026,
<https://aws.plainenglish.io/how-i-prepared-for-the-aws-data-engineer-associate-exam-and-what-id-do-differently-3425365b46d4>
32. AI/ML Starter Package - Modern Data Architecture Accelerator - AWS Documentation, accessed February 25, 2026,
<https://docs.aws.amazon.com/solutions/latest/modern-data-architecture-accelerator/ai-ml-starter-package.html>
33. Data Pipelines on AWS: 12 Services & How to Use Them Effectively | Dagster Guides, accessed February 25, 2026,
<https://dagster.io/guides/data-pipelines-on-aws-12-services-how-to-use-them-effectively>
34. Databricks vs Redshift: Lakehouse | Data Warehouse - dateonic., accessed February 25, 2026,
<https://dateonic.com/databricks-vs-amazon-redshift-comparison/>
35. Loading data from Amazon S3 - Amazon Redshift - AWS Documentation, accessed February 25, 2026,
https://docs.aws.amazon.com/redshift/latest/dg/t_Loading-data-from-S3.html
36. Python script to load data from AWS S3 to Redshift - Stack Overflow, accessed February 25, 2026,
<https://stackoverflow.com/questions/52362145/python-script-to-load-data-from-aws-s3-to-redshift>
37. Build and manage your modern data stack using dbt and AWS Glue through dbt-glue, the new “trusted” dbt adapter | AWS Big Data Blog, accessed February 25, 2026,
<https://aws.amazon.com/blogs/big-data/build-and-manage-your-modern-data-stack-using-dbt-and-aws-glue-through-dbt-glue-the-new-trusted-dbt-adapter/>
38. Redshift configurations | dbt Developer Hub, accessed February 25, 2026,
<https://docs.getdbt.com/reference/resource-configs/redshift-configs>
39. How dbt enhances your Redshift data stack, accessed February 25, 2026,
<https://www.getdbt.com/blog/redshift-dbt>
40. Databricks Lakebase. A Tale of Two Workloads | by Sumer Anand | Medium, accessed February 25, 2026,

- <https://medium.com/@sumeranand/databricks-lakebase-3ae5348e6551>
41. Databricks Lakebase is now Generally Available, accessed February 25, 2026, <https://www.databricks.com/blog/databricks-lakebase-generally-available>
 42. Decoupling Compute and Storage in Postgres: The Architectural Implications of Databricks Lakebase | by Alexey Evlampiev | Feb, 2026 | Medium, accessed February 25, 2026, <https://medium.com/@alexey.evlampiev/decoupling-compute-and-storage-in-postgres-the-architectural-implications-of-databricks-lakebase-221c503d94b8>
 43. Databricks vs Redshift: Scalability & Skills | Digiqt Blog, accessed February 25, 2026, <https://digiqt.com/blog/databricks-redshift-comparison/>
 44. Data warehousing architecture | Databricks on AWS, accessed February 25, 2026, <https://docs.databricks.com/aws/en/sql/get-started/data-warehousing-concepts>
 45. Reverse ETL with Lakebase instances | Databricks on AWS, accessed February 25, 2026, <https://docs.databricks.com/aws/en/oltp/instances/sync-data/sync-table>
 46. Databricks Lakebase: Unified OLTP on the Lakehouse Platform - Medium, accessed February 25, 2026, <https://medium.com/@sreedeeptsinha/databricks-lakebase-unified-oltp-on-the-lakehouse-platform-e2f18eeee882>
 47. How to use Lakebase as a transactional data layer for Databricks Apps, accessed February 25, 2026, <https://www.databricks.com/blog/how-use-lakebase-transactional-data-layer-databricks-apps>
 48. Announcing Lakebase Public Preview | Databricks Blog, accessed February 25, 2026, <https://www.databricks.com/blog/announcing-lakebase-public-preview>
 49. What Is a Lakebase? | Databricks Blog, accessed February 25, 2026, <https://www.databricks.com/blog/what-is-a-lakebase>
 50. Core concepts | Databricks on AWS, accessed February 25, 2026, <https://docs.databricks.com/aws/en/oltp/projects/core-concepts>
 51. Lakebase Postgres | Databricks on AWS, accessed February 25, 2026, <https://docs.databricks.com/aws/en/oltp/>
 52. Databricks Lakebase setup | dbt Developer Hub, accessed February 25, 2026, <https://docs.getdbt.com/docs/core/connect-data-platform/lakebase-setup>
 53. Navigating Modern Data Architecture: DW, Lakehouse, and Lakebase Explained, accessed February 25, 2026, <https://sqltechblog.com/2025/07/28/navigating-modern-data-architecture-dw-lakehouse-and-lakebase-explained/>
 54. Databricks Lakehouse vs Lakebase: Key Differences Explained - B EYE, accessed February 25, 2026, <https://b-eye.com/blog/databricks-lakebase-vs-lakehouse/>
 55. Databricks vs Redshift : r/dataengineering - Reddit, accessed February 25, 2026, https://www.reddit.com/r/dataengineering/comments/14s2msk/databricks_vs_redshift/
 56. Migrating from Redshift to Databricks: A Field Guide for Data Teams, accessed February 25, 2026, <https://www.databricks.com/blog/migrating-redshift-databricks-field-guide-data-teams>

57. Data Lakes vs Data Warehouses: What Your Organization Needs to Know - Databricks, accessed February 25, 2026, <https://www.databricks.com/blog/data-lakes-vs-data-warehouses-what-your-or-ganization-needs-know>
58. Databricks Lakebase GA: The Future Of AI Databases - Prolifics, accessed February 25, 2026, <https://prolifics.com/usa/resource-center/blog/databricks-lakebase-ga-explained>
59. Learn AWS online with AWS Skill Builder | Digital Training - Amazon AWS, accessed February 25, 2026, <https://aws.amazon.com/training/digital/>
60. AWS Skill Builder, accessed February 25, 2026, <https://skillbuilder.aws/>
61. Data engineer - AWS Skill Builder, accessed February 25, 2026, <https://skillbuilder.aws/category/role/data-engineer>
62. Building with Amazon Redshift Clusters - AWS Skill Builder, accessed February 25, 2026, <https://explore.skillbuilder.aws/learn/course/external/view/elearning/15681/building-with-amazon-redshift-clusters>
63. Amazon Redshift - Best Practices for Data Warehousing on AWS, accessed February 25, 2026, <https://skillbuilder.aws/learn/V6X44SFRJQ/amazon-redshift--best-practices-for-d-ata-warehousing-on-aws/DQS5GNTJ8U>
64. Request Server Error Page - AWS Skill Builder, accessed February 25, 2026, <https://skillbuilder.aws/search?searchText=managing-your-workspaces&showRedirectNotFoundBanner=true>
65. AWS Workshops, accessed February 25, 2026, <https://awsworkshop.io/>
66. Redshift - AWS Workshops, accessed February 25, 2026, <https://awsworkshop.io/tags/redshift/>
67. Modernization Workshops - AWS Workshops, accessed February 25, 2026, <https://awsworkshop.io/card/dbtlabs/>
68. Workshops | AWS Builder Center, accessed February 25, 2026, <https://builder.aws.com/build/workshops>
69. AWS Certified Data Engineer Associate Exam – DEA-C01 Study Path - Tutorials Dojo, accessed February 25, 2026, <https://tutorialsdojo.com/new-aws-certified-data-engineer-associate-dea-c01/>